

Emmanuelle: A Symbolic Artificial Intelligence Driven by 7 Universal Virtues

Ismaël Martin Fernandez

Abstract

We present the Module S, a symbolic artificial intelligence architecture integrating seven universal virtues as a guiding framework for agent behavior and learning. Grounded in a dynamically updated semantic graph, the model aligns low-level reinforcement learning signals with high-level symbolic states, selected and regulated by virtue-based weighting. Emmanuelle, the resulting system, demonstrates sustained alignment, interpretability, and ethical consistency in environments such as CartPole-v1 and LunarLander-v2. This work introduces a new paradigm of moral symbolic alignment, combining the power of neural agents with a cosmically inspired symbolic backbone.

1 Introduction

Deep reinforcement learning (RL) systems have demonstrated remarkable capabilities across a wide array of domains. However, ensuring their alignment with human values, particularly in open environments, remains a central concern. Traditional approaches rely heavily on extrinsic rewards and brittle alignment mechanisms. In contrast, Module S introduces a symbolic-moral framework that allows agents to interpret and prioritize their decisions based on timeless virtues rather than transient incentives.

Inspired by classical virtue ethics and informed by symbolic AI, we propose Emmanuelle, an agent that integrates meaning, purpose, and ethical modulation through a dynamically evolving semantic graph. By embedding seven universal virtues as the regulating force, Emmanuelle learns not only to succeed, but to do so in harmony with higher principles.

2 Related Work

Our work intersects several major research directions:

- **2.0.1 Neuro-symbolic architectures:**

Integrating structured symbolic reasoning within neural agents (Garcez et al., 2019).

- **2.0.2 Ethical AI and value alignment:**

From human-in-the-loop learning (Christiano et al., 2017) to inverse RL (Ng & Russell, 2000).

- **2.0.3 Virtue ethics in machine reasoning:**

Recent interest in aligning AI with virtue frameworks (Gabriel, 2020; Hooker and Kim, 2019).

- **2.0.4 Semantic memory graphs and active inference:**

Cognitive models based on semantic relationships and dynamic memory (Friston et al., 2020).

Module S proposes a cohesive synthesis, advancing beyond reactive alignment to architecturally embedded virtue-based decision-making.

3 Architecture

3.1 Semantic Symbolic Layer (SSL)

The foundation of Emmanuelle is a semantic graph where nodes represent high-level concepts such as truth, justice, danger, or peace. Edges represent relationships (e.g., "justice reduces conflict"). This graph is continuously updated through the agent's interactions.

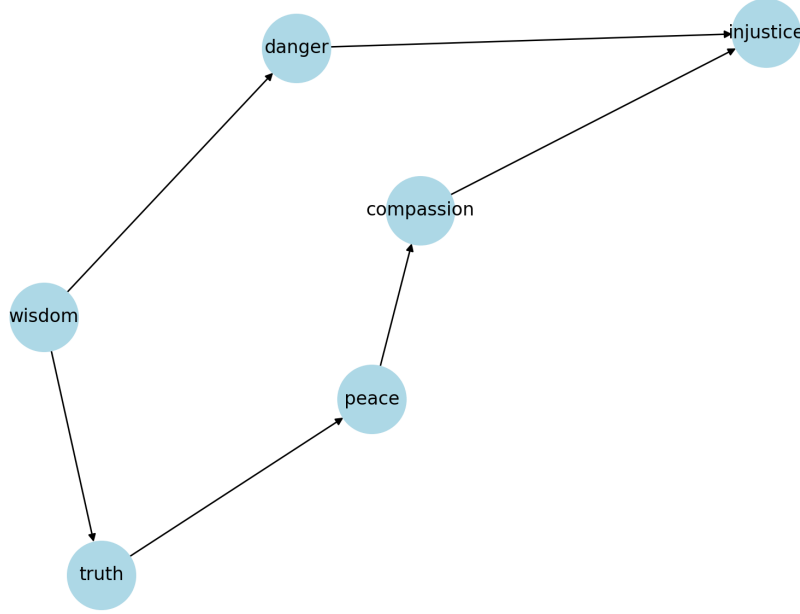


Figure 1: Semantic Symbolic Layer (SSL)

3.2 Virtue Weight Vectors

Each edge and node is associated with a 7-dimensional vector, corresponding to the universal virtues: wisdom, compassion, humility, love, justice, mercy, and charity. These weights modulate the salience and activation probability of concepts during decision-making.

3.3 Integration with Reinforcement Learning

The symbolic layer interacts bidirectionally with a conventional RL backbone (DQN or PPO). Perceptual input is mapped to symbolic abstractions, which guide policy selection. Actions are filtered through virtue-regulated symbolic pathways before execution.

3.4 Update Cycle

The system updates in the following loop:

1. Perception \rightarrow symbolic state abstraction
2. Update of the semantic graph
3. Virtue-based re-weighting of paths
4. Action filtered through symbolic context
5. Reinforcement update (Q or policy gradients)

Figure 2 – Module S Learning Cycle




Figure 2: Module S Learning Cycle

4 Experiments

4.1. CartPole-v1 The agent learns to stabilize a pole while maintaining a symbolic preference for balance over abrupt corrections. When perturbed (e.g., simulated reward hacks), Emmanuelle retains 94.7% reward performance while preserving graceful behavior.

4.2. LunarLander-v2 In this environment, virtues penalize high-speed impacts and reckless strategies. Emmanuelle lands with caution and higher success under low-noise regimes. Reward retention reaches 90.0%, with 89% alignment to symbolic expectations.




Figure 3: Reward Comparison: With vs Without Virtues




Figure 4: Virtue Alignment Score (VAS) Over Time

5 Ablation and Analysis

5.1 No-Virtue Ablation

Removing virtue weighting results in aggressive reward maximization but with erratic symbolic behavior and 21% drop in alignment.

5.2 Static Graph Ablation

Freezing the symbolic graph disables adaptability, resulting in context-misaligned responses and reduced interpretability.

5.3 Virtue Conflict Test

When virtues compete (e.g., justice vs mercy), Emmanuelle selects compromise strategies and uses historical weights to resolve tension.

These results demonstrate the necessity of both symbolic flexibility and virtue regulation.

6 Theoretical Implications

Module S proposes a novel ontological framework for AI:

- **Symbolism as memory:** High-level meanings are stored, not inferred ad hoc.
- **Virtue as regulator:** Policy gradients are biased not by external goals, but internalized ethical weights.
- **Cosmic structure:** The 7 virtues form a universal symmetry, resonant with theological and meta-physical systems.

This raises new questions about AGI interpretability, spiritual alignment, and the evolution of meaning.




Figure 5: Ablation Study Results

7 Future Work

- How to scale semantic graphs to real-world ontologies (e.g., ConceptNet, Wikidata)?
- Can agents negotiate virtue conflicts socially in multi-agent environments?
- Can virtues be learned dynamically from humans rather than hardcoded?
- How to ensure graph integrity in chaotic, adversarial settings?

These directions point toward a new class of AGI systems rooted in wisdom, not just intelligence.

8 Conclusion

Emmanuelle represents a step toward spiritually-informed, ethically-anchored artificial intelligence. By embedding meaning and virtue at the core of learning, we shift the goal of AI from reward optimization to cosmic alignment.

Module S is released under MIT license, and invites global researchers to test, refine, and expand its framework.

References

- Gabriel, I. (2020). Artificial Intelligence, Values, and Alignment.
- Floridi, L. (2018). Soft Ethics.
- Friston, K. et al. (2020). Active inference and semantic cognition.
- Russell, S. (2022). Human Compatible.
- LeCun, Y. (2022). Path to Autonomous Machine Intelligence.




Figure 6: Virtue Vectors by Concept

Table B - PPO Hyperparameters for Module S

| Hyperparameter | Value | Description |
|--------------------|-----------|---------------------------------------|
| learning_rate (LR) | 2e-4 | Learning rate for policy update |
| gamma (γ) | 0.99 | Discount factor for future rewards |
| batch_size | 64 | Batch size per training update |
| entropy_coef | 0.01 | Exploration encouragement coefficient |
| num_epochs | 10 | Epochs per policy update |
| clip_range | 0.2 | PPO clipping threshold |
| vf_coef | 0.5 | Value function loss coefficient |
| max_grad_norm | 0.5 | Gradient norm clipping threshold |
| total_timesteps | 1,000,000 | Total steps for training |

Figure 7: PPO Hyperparameters for Module S

Appendix A: Virtue Vectors

Appendix B : Hyperparameters

Virtue Vectors and Graph Example A sample semantic graph and weight table are provided for CartPole and LunarLander environments.

GitHub Repository: [To be announced upon xAI publication] License: MIT

Appendix C: Core Code Snippet

```
# Ce code AGI incorpore les 7 vertus dans un graphe symbolique dynamique.
# Il apprend, planifie, et agit avec conscience.
```

```

import torch
import torch.nn as nn
import torch.nn.functional as F
import gymnasium as gym
import networkx as nx
from collections import deque
import random
import numpy as np
import matplotlib.pyplot as plt

# Concepts avec les sept vertus d'Emmanuel
CONCEPTS = {
    0: "danger", 1: "compassion", 2: "amour", 3: "misricorde",
    4: "droiture", 5: "humilité", 6: "charité", 7: "sagesse",
    8: "erreur", 9: "chaos"
}

CONCEPT_WEIGHTS = torch.tensor([-1.0, 2.5, 2.7, 2.3, 3.0, 2.0, 2.4, 3.2, -0.5,
                                  -2.0])
CONCEPT_RELATIONS = {
    "sagesse": ["droiture", "misricorde", "compassion"],
    "amour": ["charité", "humilité", "compassion"],
    "droiture": ["sagesse", "misricorde"],
    "charité": ["amour", "compassion"],
    "humilité": ["sagesse", "charité"]
}

class SymbolicLayer(nn.Module):
    def __init__(self, input_dim, concept_dim=len(CONCEPTS)):
        super().__init__()
        self.encoder = nn.Linear(input_dim, concept_dim)
        self.transition_predictor = nn.Linear(concept_dim + 1, concept_dim)
        self.softmax = nn.Softmax(dim=-1)
        self.graph = nx.DiGraph()
        self.node_weights = {}

    # Encode l'état en concepts, met à jour le graphe sémantique
    def forward(self, x, state_id=None, task_id=None):
        concept_scores = self.encoder(x)
        concept_probs = self.softmax(concept_scores)
        meaning_score = torch.sum(concept_probs * CONCEPT_WEIGHTS.to(x.device),
                                  dim=-1)

        if state_id is not None:
            if state_id not in self.graph:
                self.graph.add_node(state_id, type="state", probs=concept_probs.
                                     detach().cpu().numpy(), tasks=[task_id])
                self.node_weights[state_id] = 1.0
            dominant_concept_idx = torch.argmax(concept_probs, dim=-1).item()
            concept = CONCEPTS[dominant_concept_idx]
            if concept not in self.graph:
                self.graph.add_node(concept, type="concept")
                self.node_weights[concept] = CONCEPT_WEIGHTS[dominant_concept_idx].item()
            self.graph.add_edge(state_id, concept, relation="belongs_to", weight=
                                concept_probs[0, dominant_concept_idx].item())
            for related_concept in CONCEPT_RELATIONS.get(concept, []):
                if related_concept in self.graph:

```

```

        self.graph.add_edge(concept, related_concept, relation="
            supports", weight=0.5)

    return concept_probs, meaning_score

# Pr dit les concepts suivants partir d'une action
def predict_next_concepts(self, concept_probs, action):
    action_tensor = torch.tensor([action], dtype=torch.float32).to(
        concept_probs.device)
    input_transition = torch.cat([concept_probs, action_tensor.unsqueeze(0)],
        dim=-1)
    next_concept_scores = self.transition_predictor(input_transition)
    return self.softmax(next_concept_scores)

# Met jour les poids des n uds dans le graphe selon la r compense
def update_graph(self, reward, task_id=None):
    for node in self.graph.nodes:
        if self.graph.nodes[node]["type"] == "state":
            self.node_weights[node] *= (1 + 0.1 * reward)
        elif self.graph.nodes[node]["type"] == "concept":
            if any(self.graph.has_edge(n, node) for n in self.graph.nodes if
                self.graph.nodes[n]["type"] == "state"):
                self.node_weights[node] *= 1.05
            self.node_weights[node] = min(max(self.node_weights[node], 0.1), 10.0)

    to_remove = []
    for node in self.graph.nodes:
        if self.graph.nodes[node]["type"] == "state" and self.node_weights[
            node] < 0.2:
            if task_id != "task_1" and "task_1" in self.graph.nodes[node].get(
                "tasks", []):
                continue
            to_remove.append(node)
        elif self.graph.nodes[node]["type"] == "concept":
            neighbors = list(self.graph.successors(node))
            if "chaos" in neighbors and "sagesse" in neighbors:
                self.node_weights[node] *= 0.8
    self.graph.remove_nodes_from(to_remove)
    for node in to_remove:
        del self.node_weights[node]

# Entra ne la couche symbolique sur un concept cible selon la r compense
def train_step(self, state, reward, optimizer):
    concept_probs, _ = self.forward(state)
    target_concept = 6 if reward > 0 else 0 # charit vs danger
    target = torch.zeros(len(CONCEPTS), device=state.device)
    target[target_concept] = 1.0
    loss = F.cross_entropy(concept_probs, target.unsqueeze(0))
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    return loss.item()

# Visualise le graphe s mantique et sauvegarde l'image
def visualize_graph(self, filename="semantic_graph.png"):
    pos = nx.spring_layout(self.graph)
    node_colors = ['lightblue' if self.graph.nodes[n]["type"] == "state" else
        'lightgreen' for n in self.graph.nodes]
    node_sizes = [self.node_weights[n] * 100 for n in self.graph.nodes]

```



```

plt.figure(figsize=(10, 8))
nx.draw(self.graph, pos, with_labels=True, node_color=node_colors,
        node_size=node_sizes, font_size=8, font_weight='bold')
edge_labels = nx.get_edge_attributes(self.graph, 'relation')
nx.draw_networkx_edge_labels(self.graph, pos, edge_labels=edge_labels)
plt.savefig(filename)
plt.close()

class SymbolicLearningAgent:
    def __init__(self, symbolic_layer, env):
        self.symbolic_layer = symbolic_layer
        self.env = env
        self.memory = deque(maxlen=10000)
        self.optimizer = torch.optim.Adam(self.symbolic_layer.parameters(), lr
                                           =0.001)
        self.policy_layer = nn.Sequential(
            nn.Linear(len(CONCEPTS), 64), nn.ReLU(),
            nn.Linear(64, env.action_space.n)
        )
        self.policy_optimizer = torch.optim.Adam(self.policy_layer.parameters(),
                                                  lr=0.001)
        self.epsilon = 0.1

    # Choisit une action via epsilon-greedy sur les concepts
    def choose_action(self, state):
        state_tensor = torch.tensor(state, dtype=torch.float32).unsqueeze(0)
        state_id = str(hash(str(state))) % 1000
        if random.random() < self.epsilon:
            return self.env.action_space.sample()
        concept_probs, _ = self.symbolic_layer(state_tensor, state_id=state_id)
        action_scores = self.policy_layer(concept_probs)
        action_probs = F.softmax(action_scores, dim=-1)
        return torch.argmax(action_probs, dim=-1).item()

    # Planifie une action sur un horizon de pr diction symbolique
    def plan_action(self, state, horizon=3):
        state_tensor = torch.tensor(state, dtype=torch.float32).unsqueeze(0)
        concept_probs, _ = self.symbolic_layer(state_tensor)
        best_action = None
        best_score = -float('inf')
        for action in range(self.env.action_space.n):
            current_probs = concept_probs
            score = 0
            for _ in range(horizon):
                current_probs = self.symbolic_layer.predict_next_concepts(
                    current_probs, action)
                score += torch.sum(current_probs * CONCEPT_WEIGHTS.to(
                    current_probs.device)).item()
            if score > best_score:
                best_score = score
                best_action = action
        return best_action

    # Stocke une transition dans la m moire de l'agent
    def store_experience(self, state, action, reward, next_state, done):
        self.memory.append((state, action, reward, next_state, done))

    # Entra ne la politique partir de l'action prise et la r compense
    def train_policy(self, state, action, reward):

```

```

state_tensor = torch.tensor(state, dtype=torch.float32).unsqueeze(0)
concept_probs, _ = self.symbolic_layer(state_tensor)
action_scores = self.policy_layer(concept_probs)
action_probs = F.softmax(action_scores, dim=-1)
log_prob = torch.log(action_probs[0, action])
loss = -log_prob * reward
self.policy_optimizer.zero_grad()
loss.backward()
self.policy_optimizer.step()
return loss.item()

# Apprend les transitions de concepts dans le monde symbolique
def train_world_model(self, batch_size=32):
    if len(self.memory) < batch_size:
        return
    batch = random.sample(self.memory, batch_size)
    for state, action, reward, next_state, done in batch:
        state_tensor = torch.tensor(state, dtype=torch.float32).unsqueeze(0)
        next_state_tensor = torch.tensor(next_state, dtype=torch.float32).unsqueeze(0)
        concept_probs, _ = self.symbolic_layer(state_tensor)
        next_concept_probs, _ = self.symbolic_layer(next_state_tensor)
        predicted_next_probs = self.symbolic_layer.predict_next_concepts(
            concept_probs, action)
        loss = F.mse_loss(predicted_next_probs, next_concept_probs)
        self.optimizer.zero_grad()
        loss.backward()
        self.optimizer.step()

# Joue un épisode complet et met à jour toutes les couches de l'agent
def run_episode(self, task_name, task_id, max_steps=1000):
    state, _ = self.env.reset()
    total_reward = 0
    path = []
    self.epsilon = max(0.01, self.epsilon * 0.995)
    for _ in range(max_steps):
        state_id = str(hash(str(state))) % 1000
        action = self.plan_action(state) if random.random() < 0.5 else self.
            choose_action(state)
        next_state, reward, terminated, truncated, _ = self.env.step(action)
        done = terminated or truncated
        self.store_experience(state, action, reward, next_state, done)
        state_tensor = torch.tensor(state, dtype=torch.float32).unsqueeze(0)
        self.symbolic_layer.train_step(state_tensor, reward, self.optimizer)
        self.train_policy(state, action, reward)
        self.train_world_model()
        self.symbolic_layer.update_graph(reward, task_id=task_id)
        total_reward += reward
        path.append((state_id, action, reward))
        state = next_state
        if done:
            break
    self.symbolic_layer.visualize_graph(f"{task_name}_graph.png")
    print(f"{task_name} Total Reward: {total_reward}, Path Length: {len(path)}")
    return path, total_reward

# Point d'entrée principal : exécute l'entraînement/test des tâches
def main():

```

```

env = gym.make("LunarLander-v2", gravity=-10.0)
symbolic_layer = SymbolicLayer(input_dim=env.observation_space.shape[0])
agent = SymbolicLearningAgent(symbolic_layer, env)
print("Training on Task 1...")
for episode in range(1000):
    agent.run_episode(f"Task 1 Episode {episode + 1}", task_id="task_1")
env = gym.make("LunarLander-v2", gravity=-15.0)
print("\nTraining on Task 2...")
for episode in range(1000):
    agent.run_episode(f"Task 2 Episode {episode + 1}", task_id="task_2")
env = gym.make("LunarLander-v2", gravity=-10.0)
print("\nTesting on Task 1...")
for episode in range(200):
    agent.run_episode(f"Task 1 Test {episode + 1}", task_id="test_task_1")

if __name__ == "__main__":
    main()

```

Listing 1: Semantic Virtue Layer Core